

DVA104

Data Structures, Algorithms, and Program Development

Gabriele Capannini

`gabriele.capannini@mdh.se`




Mälardalen University

October 1, 2018


Trees are a very common data structure used for:

- representing the file folders structure
- evaluating arithmetical expressions
- compiling a program
- performing effective search
- ...



Trees are special type of graphs so that they inherit many terms. Some new ones are listed here:


- **Root** is the top node and is unique.
- **Child** is a node directly connected to another one moving downward.
- **Parent** is the reverse relation of child.
- **Siblings** are nodes sharing the same parent.
- **Leaves** is a node having no children.
- An **empty tree** is a tree without nodes.
- Node **depth** is the number of edges in its path from the root.
- **Level** of a node is its depth + 1.
- A **tree level** is the set of nodes at the same level.
- **Degree** of a node is its number of children.



- \textcircled{A} is the root.
- \textcircled{B} , \textcircled{D} , \textcircled{E} , \textcircled{G} , \textcircled{H} are leaves.

- \textcircled{D} is child of \textcircled{C} hence \textcircled{C} is parent of \textcircled{D} .
- \textcircled{B} and \textcircled{C} are siblings.
- Depth of \textcircled{D} is 2 since its path from \textcircled{A} consists of 2 edges: $\textcircled{A} \rightarrow \textcircled{C} \rightarrow \textcircled{D}$.
- The level 3 of the tree consists of \textcircled{D} , \textcircled{E} , \textcircled{F} .
- Degree of \textcircled{C} is 3.

Example:




- **A** is the root of the tree of which nodes are: **A**, **B**, **C**, **D**, **E**, **F**, and **G**.
- **C** is the root of the **subtree** of which nodes are: **C**, **D**, **E**, **F**, and **G**.
- **D** is the root of the **subtree** of which nodes are: **D**, **F**, and **G**.
- **F** is the root of the **subtree** with only one node: **F**.

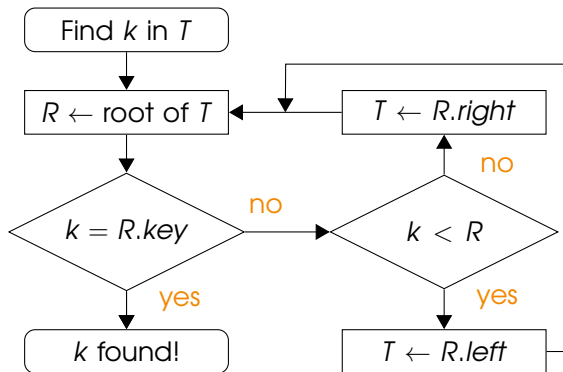
A subtree is a tree as well! 🙌

- A **binary tree** is a tree of which nodes have a degree ≤ 2 .
- So at most two subtrees can be generated by each node of a binary tree (a.k.a. **left-** and **right-subtree**).
- A **Binary Search Tree** (BST) is a Binary Tree (BT) of which nodes are associated to a value (a.k.a. **key**). Moreover, in a BST, nodes are kept sorted: let (k) be a node, all keys stored in its left-subtree are less than k , while keys that are equal or greater than k are stored in the right-subtree.


Example: In the BST on the right, all values greater than 3 and less than 7 are stored in the right-subtree of (3) and the left-subtree of (7) .



BST: how to search for a value?



Example: find 4



In the beginning we have an empty tree:




...Then we insert 5 which has no children:




...Then we insert 3 which is less than 5:



... Then we insert 8 which is greater than 5:




... Then we insert 4 which is less than 5 and greater than 3:




Building a BST (continued)

...Then we insert 9 which is greater than 5 and 8:




...Finally we insert 2 and 7:




Searching in a BST is fast!

Find 4



found!

Find 6



not found!

- In both cases, we found the answer by visiting 3 nodes while the tree contains 7 nodes.
- We also needed to visit a lot fewer nodes for entering a new item.


Does the tree T contains the value k ?

```
1: function TREEHASVALUE( $T, k$ ):  
2:   if  $T == \emptyset$  then  
3:     return NO  
4:   end if  
5:   if  $T.key == k$  then  
6:     return YES  
7:   end if  
8:   if  $k < T.key$  then  
9:     return TREEHASVALUE( $T.left, k$ )  
10:  else  
11:    return TREEHASVALUE( $T.right, k$ )  
12:  end if  
13: end function
```

Insert the value k in the tree T .



```
1: function TREECREATENode(key  $k$ ):  
2:   return node such that  $\text{key}=k$ ,  $\text{left}=\emptyset$ , and  $\text{right}=\emptyset$   
3: end function  
  
1: function TREEINSERTVALUE( $T$ ,  $k$ ):  
2:   if  $T == \emptyset$  then  
3:      $T \leftarrow \text{TREECREATENode}(k)$   
4:   end if  
5:   if  $T.\text{key} == k$  then  
6:     Raise exception  
7:   end if  
8:   if  $k < T.\text{key}$  then  
9:     TREEINSERTVALUE( $T.\text{left}$ ,  $k$ )  
10:  else  
11:    TREEINSERTVALUE( $T.\text{right}$ ,  $k$ )  
12:  end if  
13: end function
```

Is the following tree a BST?



BSTs (as well as BTs) can be easily implemented by means of an array:

- root is registered in the first position, i.e., 0
- left child of the node stored in position i is in position $2i + 1$
- right child of the node stored in position i is in position $2i + 2$




Cons:

- When array is full you need to resize it which can be costly.
- Array has 'holes' if the tree is not complete.

Definition: a binary tree is **complete** if all levels are completely filled except possibly the last level that has all keys as left as possible.

Example:



BTs can be also implemented by means of a LL.



Example of a tree node with integer keys:

```
struct btreenode {  
    int key;  
    treeNode* left;  
    treeNode* right;  
};  
typedef struct btreenode BTreeNode;
```

When a new node is created `left` and `right` are `NULL` which means that, initially, a new node has no children:

```
BTreeNode* createNewBTreeNode(const int key) {  
    BTreeNode* newnode = (BTreeNode*)malloc(sizeof(BTreeNode));  
    newnode->left = newnode->right = NULL;  
    newnode->key = key;  
    return newnode;  
}
```

Example:



∅ = NULL

- Each pointer points to a whole `BTNode` not only to `key`.
- Memory management is simplified but, now, each node needs more memory to be stored.

Just like a LL or an array, we may want to visit every node in a tree.

However, depending on our purpose, it is not obvious in which order you want to visit the nodes in a tree.

For example:

- Should the nodes on top be visited before the bottom?
- Should we visit left-hand tree trees before right-hand tree trees?

There are many (recursive) alternatives:

Pre-order (LR):

1. visit the current node,
2. visit the left-subtree,
3. visit the right-subtree.

In-order (LR):

1. visit the left-subtree,
2. visit the current node,
3. visit the right-subtree.

Post-order (LR):

1. visit the left-subtree,
2. visit the right-subtree,
3. visit the current node.

Pre-order (RL):

1. visit the current node,
2. visit the right-subtree,
3. visit the left-subtree.

In-order (RL):

1. visit the right-subtree,
2. visit the current node,
3. visit the left-subtree.


Post-order (RL):

1. visit the right-subtree,
2. visit the left-subtree,
3. visit the current node.

Tree traversal (example)

We want to print the following tree in Pre-order (LR) manner:

```
1: procedure VISIT(T)
2:   PRINT(T.key)
3:   VISIT(T.left)
4:   VISIT(T.right)
5: end procedure
```




*dashed orange boxes enclose trees on which a visit procedure is active.

Printed nodes: 5

Tree traversal (example)

We want to print the following tree in Pre-order (LR) manner:

```
1: procedure VISIT(T)
2:   PRINT(T.key)
3:   VISIT(T.left)
4:   VISIT(T.right)
5: end procedure
```




*dashed orange boxes enclose trees on which a visit procedure is active.

Printed nodes: 5, 3

Tree traversal (example)

We want to print the following tree in Pre-order (LR) manner:

```
1: procedure VISIT(T)
2:   PRINT(T.key)
3:   VISIT(T.left)
4:   VISIT(T.right)
5: end procedure
```




*dashed orange boxes enclose trees on which a visit procedure is active.

Printed nodes: 5, 3, 2

Tree traversal (example)

We want to print the following tree in Pre-order (LR) manner:

```
1: procedure VISIT(T)
2:   PRINT(T.key)
3:   VISIT(T.left)
4:   VISIT(T.right)
5: end procedure
```




*dashed orange boxes enclose trees on which a visit procedure is active.

Printed nodes: 5, 3, 2, 4

Tree traversal (example)

We want to print the following tree in Pre-order (LR) manner:

```
1: procedure VISIT(T)
2:   PRINT(T.key)
3:   VISIT(T.left)
4:   VISIT(T.right)
5: end procedure
```




*dashed orange boxes enclose trees on which a visit procedure is active.

Printed nodes: 5, 3, 2, 4, 8

Tree traversal (example)

We want to print the following tree in Pre-order (LR) manner:


```
1: procedure VISIT(T)
2:   PRINT(T.key)
3:   VISIT(T.left)
4:   VISIT(T.right)
5: end procedure
```



*dashed orange boxes enclose trees on which a visit procedure is active.

Printed nodes: 5, 3, 2, 4, 8

What's the result in this case? Which kind of traversal are we performing?



```
void printBST(BTNode* tree) {  
    if(!root)  
        return;  
    printBST(tree->left);  
    printf("%d\n", tree->key);  
    printBST(tree->right);  
}
```


Adding a node into a BST is 'easy'...

While removing can lead to different cases:

- remove a leaf,
- remove a node with a child,
- remove a node with two children.

Removing a leaf

```
typedef struct breenode BTreeNode;  
typedef BTreeNode* BTree;  
void remove(BTree* tree /*double pointer*/, int key) {  
    ...  
}  
BTree root;  
...  
remove(&root, 9); //initially tree points to root
```




first of all we traverse
the tree until tree
point to the pointer to
⑨ ...

Removing a leaf (continued)

We recursively set `tree` with the address of the pointers belonging to the path from ⑤ to ⑨. For example, the first recursive call will look like:

```
BTNode* node = *tree;
...
if(key<node->key) // i.e. 9<5 is FALSE, go right
...
else
    remove(&node->right, 9); // first recursive call
```




in the second step,
tree points to the
right-subtree of ⑤

Removing a leaf (continued)

We reached ⑨ and `tree` points to the right-subtree of ⑧ :

```
BTNode* node = *tree;  
if(key==node->key) // i.e. 9==9 is TRUE  
{  
    free(*tree); //free memory  
    *tree = NULL; //update pointer  
}
```




after removing ⑨ we
get a tree like this

Removing nodes with one child

This time we want to remove a node like ⑧ :

```
BTree root;  
...  
remove(&root, 8); //initially tree points to root
```




first of all we traverse
the tree until we
reach ⑧ ...

Removing nodes with one child (continued)




As in the previous case, we need to reach the node of which child points to the target value: in this case it is 8. In this case, however, (8) has a child:



Removing nodes with one child (continued)



We replace the child-pointer pointed by `tree` with the child of the node to remove:



That is:


```
BTNode* tmp = *tree;  
*tree = *tree->right;  
free(tmp);
```

Notice that 'removing a leaf' is just a 'simplified' case of this since in that case `*tree->right` is `NULL`.

Removing nodes with two children

This time we want to remove a node like ⑤ :

```
BTree root;  
...  
remove(&root, 5); //initially tree points to root
```



In this case we need to find a good candidate to replace ⑤ : who?



Hint: remember the definition of BST!

Removing nodes with two children (continued)

Removing ⑤ :

we can alternatively choose:

- the **max** value stored in the **left-subtree**, that is 4
- the **min** value stored in the **right-subtree**, that is 7





By the way, how to find the min and max value in a BST?

Removing nodes with two children (continued)

Removing ⑤ :

Once the node has been replaced, we proceed by removing the copied one:




Note that removing such a node must fall in one of the two previous cases: 'removing a leaf' or 'removing a node with one child'.

Questions:


- How does a BST look like if we insert the sequence 1,2,3,4,5?
- How many nodes are there in the tree?
- Is it still quick to search in this tree?
- Is it faster than a linked list?
- How many levels have the tree?

Answers:

- It looks like the next tree.
- There are 5 nodes.
- Searching is not faster than a LL.
- The tree has 5 levels.
- Such a tree is called unbalanced and it is not effective in searching.
- If the number of nodes equals the number of levels then the tree is a LL! (degenerate tree)



So, what is the min number of levels for a tree with 5 nodes?




- The capacity of a BST with 1 level is 1 (just the root)
- The capacity of a BST with 2 levels is 3 (root+2ch, i.e., 1+2)
- The capacity of a BST with 3 levels is 7 (i.e., 3+4)
- The capacity of a BST with 4 levels is 15 (i.e., 7+8)

In general, the capacity of a BST with L levels is $2^L - 1$

Hence the minimum number of levels for a BST with 5 nodes is 3!

For example it could be:



- Is that BST sorted? Yes
- How many nodes? 5
- How many levels? 3
- Is it effective for searching? Yes, max 3 steps

If we can not store more than $2^L - 1$ nodes in a BST with L levels, how many levels do we need to store n nodes?

$$n \leq 2^L - 1$$

$$n + 1 \leq 2^L$$

$$\log_2(n + 1) \leq \log_2(2^L) = L$$

$$L \geq \log_2(n + 1)$$

Answer: we need at least $\log_2(n + 1)$.

Since $\log_2(n + 1)$ may contain decimals and L should be an integer, we must round up the result: $\lceil \log_2(n + 1) \rceil$.

In C, round-up can be calculated using the `ceil` function:
<http://www.cplusplus.com/reference/cmath/ceil/>

- A BST with n nodes on $\lceil \log_2(n+1) \rceil$ levels is said to be **balanced** (definition comes later).
- In the worst case, we need to visit $\lceil \log_2(n+1) \rceil$ nodes to find a given value, while in a LL this requires to visit n nodes:


For example, assume that we have stored about one million values, e.g., $n = 1048576 = 2^{20}$.

LL in the worst case, we need to test one million (i.e., 1048576) nodes to find a specific one.

BST we need to test $\lceil \log_2(n+1) \rceil$ that is 20 nodes (at most).

Tree balancing (continued)

The following plot shows the number of *visited* nodes w.r.t. the number n of nodes stored in the data structure both for the LL and the balanced BST cases:



Application:



Do you remember ADTs Set?

```
addElement ()  
removeElement ()  
isInSet ()
```

If we implemented this as a sorted balanced BST instead of an Array or LL then the `isInSet ()` function would be much faster!

Tree balancing (continued)

Definition: **balanced tree** means that the difference between the height of left- and right-subtree of each node is max 1, where the **height** of a (sub)tree is defined as the number of edges on the longest path between the root and a leaf.



According to the definition only the left one is balanced. Why?

Now, the question is: how do we get a tree balanced?

There are several variants of self-balanced BSTs:

- AVL trees
- Red/Black Trees

Otherwise we can:

- First, enter all values from the tree into a sorted array (in-order LR visit).
- Then, balance the tree (rebuild the tree with the same values).


Assume we have a sorted array with the values to store in the tree, how can we build it balanced?

For example, we want to store the values from 1 to 10 in a BST.




- How can we build it balanced?
- Which value should be the root?

Building a balanced BST (continued)




The array has 10 entries (from position 0 to 9), '5' (in position 4) is almost in the middle: $4 = \lfloor (0 + 9)/2 \rfloor$.




Now, what should be the root of the right-subtree of 5?


Building a balanced BST (continued)



The right-subtree has 5 entries (from position 5 to 9), '8' (in position 7) is in the middle: $7 = \lfloor (5 + 9)/2 \rfloor$.



We follow the same schema for adding '9' and '10' to the tree:



... Even for the left-subtree. For example, what should be the root of the left-subtree of **8**? The left-subtree covers the array positions from 5 to 6, hence we select the value in position 5 (since $\lfloor (5 + 6)/2 \rfloor = 5$).

We are building a balanced BST recursively. At each step:

- First, we choose the root from the given (sub)array.
- Then, we recursively apply the same procedure on the subarrays on the left and on the right of the chosen element for the left- and the right-subtree, respectively.

function *BUILDBST*(*node*, *array*, *first*, *last*):

...

$m \leftarrow \lfloor (first + last) / 2 \rfloor$

node.key \leftarrow *array*[*m*]

BUILDBST(*node.left*, *array*, *first*, *m* - 1)

BUILDBST(*node.right*, *array*, *m* + 1, *last*)

end function

What is the 'base-case' for this recursive function?

When '10' is added, *first* and *last* equal 9 so that even *m* is equal to 9. Hence , the recursive calls look like:

```
function BUILDBST(node, array, first=9, last=9):
```

```
...
```

```
   $m \leftarrow \lfloor (first + last)/2 \rfloor = \lfloor (9 + 9)/2 \rfloor = 9$ 
```

```
  node.key  $\leftarrow$  array[m]
```

```
  BUILDBST(node.left, array, first, m-1)
```

```
  BUILDBST(node.right, array, m+1, last)
```

```
end function
```

The function is called twice on empty portions of *array* since $first > m-1$ and $m+1 > last$, which makes no sense. Hence the function should preliminary check the values of *first* and *last*:


```
function BUILDBST(node, array, first, last):
```

```
  If  $first > last$  then exit
```

```
...
```

```
end function
```

Look at the following tree:



- Is it a BT?
- Is it a BST?
- Is it balanced?

The end... Questions?